

# Naive Bayes

A set of supervised learning algorithms based on applying Bayes theorem with the ‘naive’ assumption of independence between every pair of features.

$$P(y | x_1, \dots, x_n) = \frac{P(y)P(x_1, \dots, x_n | y)}{P(x_1, \dots, x_n)}$$

- Works well in document classification and spam filtering
- Require a small amount of training data to estimate the necessary parameters
- Fast
- Good classifier, bad estimator

## Gaussian Naive Bayes

$$P(x_i | y) = \frac{1}{\sqrt{2\pi\sigma_y^2}} \exp\left(-\frac{(x_i - \mu_y)^2}{2\sigma_y^2}\right)$$

```
>>> from sklearn import datasets
>>> iris = datasets.load_iris()
>>> from sklearn.naive_bayes import GaussianNB
>>> gnb = GaussianNB()
>>> y_pred = gnb.fit(iris.data,
iris.target).predict(iris.data)
>>> print("Number of mislabeled points out
of a total %d points :
...      %d" % (iris.data.shape[0],
(iris.target != y_pred).sum()))
Number of mislabeled points out of a total
150 points : 6
```

# Multinomial Naive Bayes

MultinomialNB implements the naive Bayes algorithm for multinomially distributed data.

- Used in text classification
- Used for discrete counts

## Bernoulli Naive Bayes

BernoulliNB implements the naive Bayes training and classification algorithms for data that is distributed according to multivariate Bernoulli distributions

- May be multiple features but each one is assumed to be a binary-valued (Bernoulli, boolean) variable
- Requires samples to be represented as binary-valued feature vectors

$$P(x_i | y) = P(i | y)x_i + (1 - P(i | y))(1 - x_i)$$

### Differs from multinomial NB's rule:

- Explicitly penalizes the non-occurrence of a feature  $i$  that is an indicator for class  $y$
- Multinomial variant would simply ignore a non-occurring feature
- Perform better on datasets with shorter documents
- Evaluate both models if time permits

## Out-of-core Naive Bayes Model Fitting

MultinomialNB, BernoulliNB, and GaussianNB expose a `partial_fit` method that can be used incrementally as done with other classifiers.

All naive Bayes classifiers support *sample weighting*.

Contrary to the `fit` method, the first call to `partial_fit` needs to be passed the list of all the expected class labels.

## Calculating Accuracy Score

```
def NBAccuracy(features_train, labels_train, features_test, labels_test):  
    from sklearn.naive_bayes import GaussianNB  
    from sklearn.metrics import accuracy_score  
  
    clf = GaussianNB()  
    clf.fit(features_train, labels_train)  
  
    pred = clf.predict(features_test)  
  
    accuracy = accuracy_score(labels_test, pred)  
    return accuracy
```

## Bayes Rule (also called Bayes Theorem)

$$P(A | B) = \frac{P(B | A) P(A)}{P(B)}$$

## Why Naive Bayes naive?

Naive bayes ignores words order, and it just cares about word frequency.